

Operating System Lab

Lab #8

Semaphores

Goal:

This lab discusses the semaphore and its application. The expectation is to develop deeper understanding of how to use a semaphore.

Preparation:

Prepare the working environment

- Check out the repo from Github classroom
 - o Use `git clone {repo URL}`
- Rebuild Docker image and run the container
 - o In a Windows Powershell or MacOS/Linux terminal, change the current folder to the “lab” folder (the one you check out).
`cd {lab_folder}`
 - o For **Mac** or **Linux**
 - You can execute the shell script “`clean_start.sh`” if you want to remove old image and create and run a new one.
`./clean_start.sh`
 - Otherwise, you can just run the following “docker run” command to run a container.
 - o For **Windows**
 - You need the following docker commands to run the container from the existing image, *ubuntu_os*.

```
docker run -itd -p8888:8888 --rm --name ubuntu-os-con -v"$PWD/share":/home/app
ubuntu_os
```

- If you do not have the image *ubuntu_os*, you need to rebuild the image

```
docker build -t ubuntu_os .
```

- o To enter the “Bash Shell” of the Ubuntu container for compiling and testing codes

```
docker exec -it ubuntu-os-con /bin/bash
```

- Use VSCode or any text editors for development
 - o In the host OS (Windows or MacOS), use terminal to change to the host’s lab folder “{check out repo folder}/share” or VSCode to open that host’s lab folder. Files you created or modified in the “{check out repo folder}/share” will be shown in docker container’s “/home/app” folder.
 - o Copy required files from “{check out repo folder}/share” to “{check out repo folder}/answer” and check in for submission.

Part I: Learn POSIX Semaphore

1. Basic

The POSIX system in Linux has a built-in semaphore library. The main functions are shown as follows.

- `sem_init()` - to initialize a semaphore
- `sem_wait()` - to wait on a semaphore
- `sem_post()` - to increment the value of a semaphore (or if it is a binary semaphore, it will release the semaphore)
- `sem_getvalue()` - to get the current value of the semaphore
- `sem_destroy()` - to destroy the semaphore

A semaphore is an object with an integer value that we usually manipulate it with `sem_wait()` and `sem_post()`. To use it, we must include the `semaphore.h`. We can initialize semaphores as follow:

```
#include <semaphore.h>
sem_t s;
sem_init(&s, 0, 1);
```

We declare a semaphore `s` and initialize it. From the Linux man page: https://man7.org/linux/man-pages/man3/sem_init.3.html, we know that the setting 0 for the second argument of `sem_init()` indicates that semaphore is shared between threads in the same process. The third argument can be any value. In above example, we pass 1 in as the third argument. After a semaphore is initialized, we can call one of two functions to interact with it, `sem_wait()` or `sem_post()`. The behavior of these two functions are seen in Figure 1.

```
1  int sem_wait(sem_t *s) {
2      decrement the value of semaphore s by one
3      wait if value of semaphore s is negative
4  }
5
6  int sem_post(sem_t *s) {
7      increment the value of semaphore s by one
8      if there are one or more threads waiting, wake one
9  }
```

Figure 1: Semaphore: definitions of `sem_wait` and `sem_post`

We can see that `sem_wait()` will either return right away (because the value of the semaphore was one or higher when we called `sem_wait()`), or it will cause the caller to suspend execution waiting for a subsequent post. Of course, multiple calling threads may call into `sem_wait()`, and thus all be queued waiting to be woken.

We can see that `sem_post()` does not wait for some particular condition to hold like `sem_wait()` does. Rather, it simply increments the value of the semaphore and then, if there is a thread waiting to be woken, wakes one of them up.

2. Binary Semaphores (Locks)

We are now ready to use a semaphore. We utilize binary semaphore as a lock shown in Figure 2.

```
sem_t m;  
sem_init(&m, 0, X); //initialize to X; what value should X be?  
  
sem_wait(&m);  
// critical section here  
sem_post(&m);
```

Figure 2: A binary semaphore implements a lock

You will see that we simply surround the critical section of interest with a `sem_wait()` and `sem_post()` pair.

To-do (Part1):

1. Suppose the value of the current semaphore is -3. How many threads are waiting to enter the critical section? Explain it in the “*answer.md*” under the item “Part1-1”?
2. Take a look at Figure 2 and answer the following questions: Determine what should be X as the initial value of the semaphore m? Explain it in the “*answer.md*” under the item “Part1-2”?
3. Check the unfinished “*binary.c*” code from your answer folder and complete the TODO sections. Answer what the functionality the code has in the “*answer.md*” under the item “Part1-3”?

```
Zem_t s;  
  
//@TODO: add necessary codes to fix this  
void *child(void *arg) {  
    sleep(4);  
    printf("child\n");  
    return NULL;  
}  
  
int main(int argc, char *argv[]) {  
    Zem_init(&s, 0);  
    printf("parent: begin\n");  
    pthread_t c;  
    printf("parent: end\n");  
    return 0;  
}
```

Figure 3: A parent waiting for its child

4. Semaphores are also useful to order events in a concurrent program. Complete a simple program called “*zemaphore.c*” shown in Figure 3. To do that, check the skeleton code “*zemaphore.c*” in the `answer` folder and **use it in your program**. Imagine a thread creates

another thread and then wants to wait for it to complete its execution. The following is the output of the program.

```
parent: begin
child
parent: end
```

Part II: The Producer/Consumer (Bounded Buffer) Problem

Our first attempt at solving the problem introduces two semaphores, **empty** and **full**, which the threads will use to indicate when a buffer entry has been emptied or filled, respectively. The code for the put and get functions is shown in Figure 4. The `put()` and `get()` are convenient functions to use buffer. Our attempted solution for the producer and consumer problem is shown in Figure 5.

```
int *buffer;
int use = 0;
int fill = 0;

void put(int value) {
    buffer[fill] = value;
    fill++;
    if (fill == max)
        fill = 0;
}

int get() {
    int tmp = buffer[use];
    use++;
    if (use == max)
        use = 0;
    return tmp;
}
```

Figure 4: The `put()` and `get()`

In this example, the producer first waits for a buffer to become empty to put data into it, and the consumer similarly waits for a buffer to become filled before using it. Let us first imagine that `max=1` (there is only one buffer in the array) and see if this works.

Consider two threads, a producer and a consumer, and a single CPU. Assume the consumer gets to run first. Thus, the consumer will hit Line C1 in Figure 5 calling `sem_wait(&full)`. Because **full** was initialized to the value 0, the call will decrement **full** (to -1), block the consumer, and wait for another thread to call `sem_post()` on **full**, as desired.

Look at the Line P4. The reason to add `put (-1)` is to make the variable “tmp” to receive -1 and then exit the while loop.

```
sem_t empty;
sem_t full;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&empty); // Line P1
        put(i);           // Line P2
        sem_post(&full);  // Line P3
    }
    put(-1);             // Line P4
    return NULL;
}

void *consumer(void *arg) {
    int i, tmp = 0;
    while (tmp != -1) {
        sem_wait(&full); // Line C1
        tmp = get();     // Line C2
        sem_post(&empty); // Line C3
        printf("%lld %d\n", (long long int) arg, tmp);
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    //...
    sem_init(&empty, 0, max); // max are empty
    sem_init(&full, 0, 0);    // 0 are full
    //...
}
```

Figure 5: Adding the full and empty conditions

To-do (Part2):

1. Imagine `max = 10` for the code shown in Figure 5. Consider multiple producers and consumers. Do we have the race condition? Please put your explanation in “[answer.md](#)” under the item “**Part2-1**”. The entire program is the “[pc_fig.c](#)” in the answer folder. You can view it if necessary. **Please also fix it if you think the program has issues.** If you fix the program, do not need to change the filename.
2. Please start from the skeleton “[pc_works_skeleton.c](#)” and write a program “[pc_works.c](#)”. “[pc_works.c](#)” is implemented learning from “[pc_fig.c](#)” but use the functions in “[common_threads.h](#)”. Therefore, in the “[pc_works.c](#)”, unlike using Linux’s semaphore used in “[pc_fig.c](#)”, you need to read the corresponding function in the “[common_thread.h](#)” to finish your implementation.